

CSCI485HW2

February 23, 2026

```
[1]: from sklearn.datasets import load_diabetes
import matplotlib
import pandas as pd
import numpy as np
import altair as alt
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn import set_config
from sklearn.metrics import r2_score
from sklearn.feature_selection import RFE
from sklearn.model_selection import cross_val_score, train_test_split
import matplotlib.pyplot as plt
from sklearn.datasets import make_regression
```

1 Task 1: Dataset Exploration

```
[2]: # Load data

diabetes = load_diabetes()

df = pd.DataFrame(data=diabetes.data, columns=diabetes.feature_names)
y = diabetes.target

df.head()
```

```
[2]:      age      sex      bmi      bp      s1      s2      s3 \
0  0.038076  0.050680  0.061696  0.021872 -0.044223 -0.034821 -0.043401
1 -0.001882 -0.044642 -0.051474 -0.026328 -0.008449 -0.019163  0.074412
2  0.085299  0.050680  0.044451 -0.005670 -0.045599 -0.034194 -0.032356
3 -0.089063 -0.044642 -0.011595 -0.036656  0.012191  0.024991 -0.036038
4  0.005383 -0.044642 -0.036385  0.021872  0.003935  0.015596  0.008142

      s4      s5      s6
0 -0.002592  0.019907 -0.017646
1 -0.039493 -0.068332 -0.092204
```

```
2 -0.002592  0.002861 -0.025930
3  0.034309  0.022688 -0.009362
4 -0.002592 -0.031988 -0.046641
```

```
[ ]: # Explore the dataset and describe features and target variables
      print(df.info())
[ ]: print(df.describe())
```

We have 10 features that will be used to predict our target:

Features:

```
age: Age of the patient.
sex: Gender of the patient.
bmi: Body mass index.
bp: Average blood pressure.
Six other quantitative variables derived from blood samples.
```

Target: A continuous variable measuring the disease progression within a one-year period. Larger values generally correspond to more severe progression of diabetes.

All our values are floats, with no nulls.

442 observations (diabetes patients).

```
[3]: # split the dataset into testing and training using an 80-20 split.

X = pd.DataFrame(data=diabetes.data, columns=diabetes.feature_names)
y = diabetes.target

# separating by X and y values will help calculate Adjusted R squared later
X_train, X_test, y_train, y_test = train_test_split(
    X, y, train_size=0.80, random_state = 42
)
```

2 Task 2: Linear Regression Model

```
[4]: # Full Model
full_model = LinearRegression()
full_model.fit(X_train, y_train)
full_pred = full_model.predict(X_test)

# Metrics for Full Model
full_mse = mean_squared_error(y_test, full_pred)
full_r2 = r2_score(y_test, full_pred)
print("Full Model Metrics:")
print(f"MSE: {full_mse:.10f}, R²: {full_r2:.10f}")
```

Full Model Metrics:

MSE: 2900.1936284935, R²: 0.4526027630

45.26% of the variation in our target is accounted for by our model.

3 Tasks 3/4: Implement Recursive Feature Elimination (RFE) & Analyze Feature Importance

```
[5]: def rfe_starting_with_all_features():
    rfe_results = []
    feature_names = diabetes.feature_names
    threshold = 0.01

    for k in range(len(feature_names), 0, -1):
        rfe = RFE(estimator=full_model, n_features_to_select=k)
        rfe.fit(X_train, y_train)

        current_r2 = r2_score(y_test, rfe.predict(X_test))

        # put 0 if the feature isn't selected
        coef_map = {name: 0.0 for name in feature_names}
        selected_features = [feature_names[i] for i in
        ↪range(len(feature_names)) if rfe.support_[i]]

        for name, coef in zip(selected_features, rfe.estimator_.coef_):
            coef_map[name] = coef

        # Store results
        coef_map['n_features'] = k
        coef_map['R2_Score'] = current_r2
        rfe_results.append(coef_map)

    table_df = pd.DataFrame(rfe_results).set_index('n_features').sort_index()

    # getting the largest count where the improvement is above the defined
    ↪threshold 0.01
    optimal_n = 1
    for i in range(2, len(feature_names), 1):
        improvement = table_df.loc[i, 'R2_Score'] - table_df.loc[i-1, ↪
        'R2_Score']
        if improvement > threshold:
            optimal_n = i
    print(f"Optimal number of features given threshold {threshold}:"
    ↪{optimal_n}")
    return table_df, optimal_n
```

```

results_table, optimal_n = rfe_starting_with_all_features()
display(results_table)

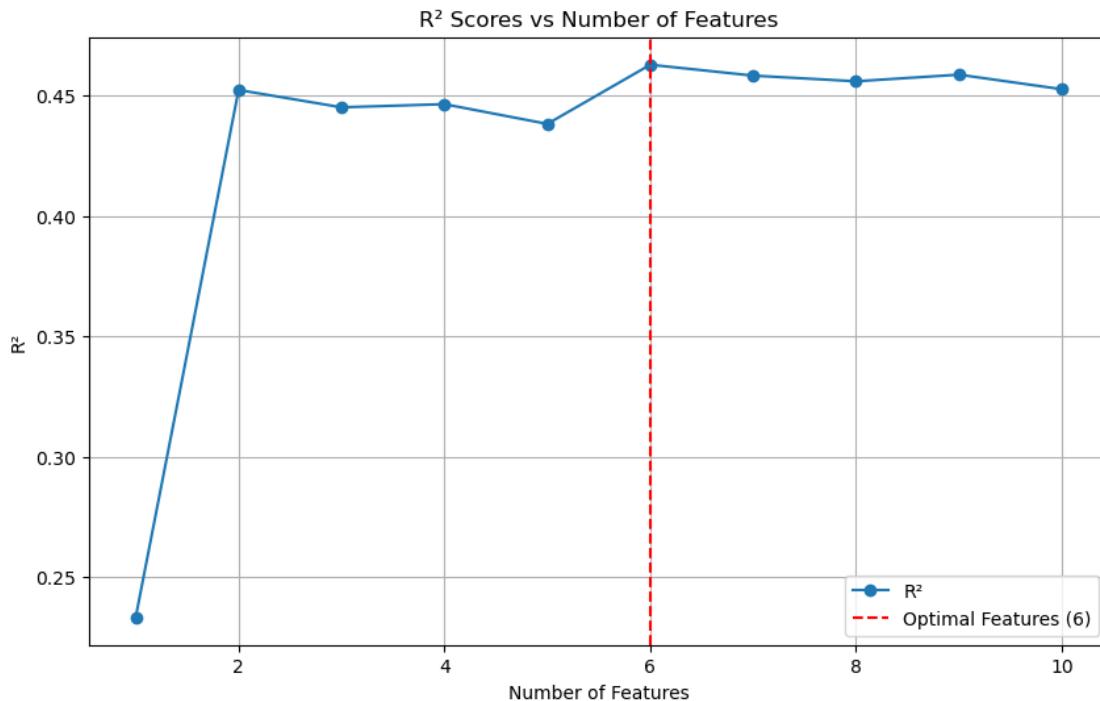
plt.figure(figsize=(10, 6))
plt.plot(results_table.index, results_table['R2_Score'], marker='o', label='R2')
plt.axvline(optimal_n, color='red', linestyle='--', label=f'Optimal Features\n→({optimal_n})')
plt.title("R2 Scores vs Number of Features")
plt.xlabel("Number of Features")
plt.ylabel("R2")
plt.legend()
plt.grid()
plt.show()

```

Optimal number of features given threshold 0.01: 6

n_features	age	sex	bmi	bp	s1	\
1	0.000000	0.000000	998.577689	0.000000	0.000000	
2	0.000000	0.000000	732.109021	0.000000	0.000000	
3	0.000000	0.000000	737.685594	0.000000	-228.339889	
4	0.000000	0.000000	691.460102	0.000000	-592.977874	
5	0.000000	0.000000	597.892739	306.647913	-655.560612	
6	0.000000	-215.267423	557.314167	350.178667	-851.515734	
7	0.000000	-235.364224	551.866448	362.356114	-660.643160	
8	0.000000	-233.754686	550.744365	363.791753	-947.823133	
9	0.000000	-236.649588	542.799508	354.211438	-936.350589	
10	37.904021	-241.964362	542.428759	347.703844	-931.488846	
n_features	s2	s3	s4	s5	s6	\
1	0.000000	0.000000	0.000000	0.000000	0.000000	
2	0.000000	0.000000	0.000000	562.226535	0.000000	
3	0.000000	0.000000	0.000000	680.224653	0.000000	
4	362.950323	0.000000	0.000000	783.168538	0.000000	
5	409.622184	0.000000	0.000000	728.643647	0.000000	
6	591.093315	0.000000	0.000000	803.121285	0.000000	
7	343.348089	0.000000	185.140764	664.774591	0.000000	
8	541.585796	172.250588	277.741072	761.921177	0.000000	
9	528.796592	167.800414	270.396514	744.447429	53.350483	
10	518.062277	163.419983	275.317902	736.198859	48.670657	
n_features	R2_Score					
1	0.233350					
2	0.452293					
3	0.445095					
4	0.446404					

5	0.438201
6	0.462777
7	0.458255
8	0.455901
9	0.458659
10	0.452603



My code above starts with 10 features and decreases by 1 feature incrementally, tracking each R^2 value and coefficients in a table. Then to find the optimal amount of features, it checks each R^2 , and if it improves by at least the desired threshold I defined to be 0.01, then it is considered the optimal model. This method allowed me to find the optimal number of features as 6 (sex, bmi, s1, s2, , s5, and s6) given a threshold of 0.01. Finally, I graphed my results.

2. The three most important features are bmi, s1, and s5 because their coefficients are furthest from 0 during each iteration, shown by my table. Those 3 features alone are able to explain 44.51% of the variance in the target. This is not far from the 46.27% of variance explained from the optimal 6-feature model. Choosing these 3 features, given they have similar R^2 , would give us a less complex model for similar model results. Changing our threshold value based on how much accuracy we're willing to give up for less complexity is the clear goal of RFE.
3. Our baseline model, with all 10 features, has R^2 of 0.4526, meaning 45.26% of the variation in our target is accounted for by our model. The RFE process removed the least important feature, or the feature with coefficient closest to 0, one at a time. My results showed that removing age, s6, s3, then s4, left us with the best r^2 given our threshold. Therefore, our optimal model explains 46.28% of the variation in the target using the 6 remaining features

sex, bmi, s1, s2, , s5, and s6.

4 Task 5: Reflection

1. I learned that the way RFE works is by iteratively eliminating the least important features to find the model that gives us the best R^2 . I learned that it's important to choose a good threshold for how much accuracy you're willing to part with in exchange for model simplicity. The difference in R^2 from a 2-feature model to our optimal model (6 features) is very small, so it may be worth using just 2 features depending on the circumstance.
2. RFE removes the coefficients with the smallest absolute value, whereas Lasso's methodology involves minimizing the sum of the absolute values of the coefficients by penalizing large coefficients. They both aim to improve model performance by reducing the number of features.
3. Age of the patient is the worst predictor (out of our 10 features) for predicting our target, followed by s6 and s3. The best predictors we have are bmi of the patient, then s5 and s1.